

Lecture 3

Dimensionality Reduction

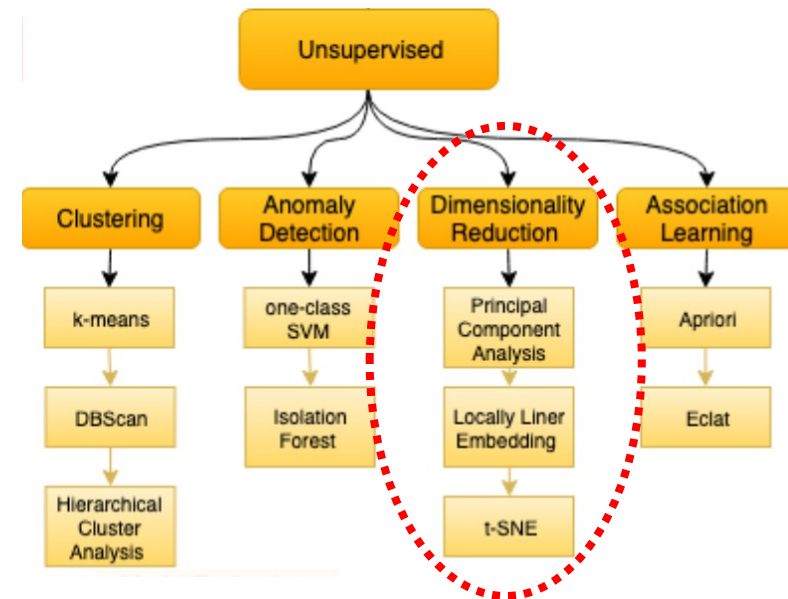
Ioanna Miliou, PhD

Senior Lecturer, Stockholm University

Unsupervised learning

Experience: objects for which **no class labels** have been given

Performance: typically concerns the ability to output useful **characterizations** (or groupings) of objects




Dimensions

Features (attributes)						Class label	
Examples (observations)	Email	All caps	No. excl. marks	Missing date	No. digits in From:	Image fraction	Spam
	e1	yes	0	no	3	0	yes
	e2	yes	3	no	0	0.2	yes
	e3	no	0	no	0	1	no
	e4	no	4	yes	4	0.5	yes
	e5	yes	0	yes	2	0	no
	e6	no	0	no	0	0	no


Data Dimensionality

- From a theoretical point of view, increasing the number of features should lead to better performance
- In practice, the inclusion of more features leads to worse performance (i.e., the **curse of dimensionality**)
- Need an **exponential** number of training examples as dimensionality increases
- Index structures fail as the dimensionality of the data increases




The Curse of Dimensionality


- Data in only one dimension is relatively **packed**
- Adding a dimension “stretches” the points across that dimension, making them **further apart**
- Adding more dimensions will make the points further apart—high dimensional data is **extremely sparse**
- **Distance measures** become meaningless



(a) 11 Objects in One Unit Bin



(b) 6 Objects in One Unit Bin



(c) 4 Objects in One Unit Bin

Data Dimensionality

- Significant improvements can be achieved by first mapping the data into a **lower-dimensional** space:


$$x = \begin{bmatrix} a_1 \\ a_2 \\ \dots \\ a_N \end{bmatrix} \longrightarrow \text{reduce dimensionality} \longrightarrow y = \begin{bmatrix} b_1 \\ b_2 \\ \dots \\ b_K \end{bmatrix} \quad (K \ll N)$$

- Dimensionality can be reduced by:
 - Combining features (**linearly** or **non-linearly**)
 - Selecting a subset of features (i.e., **feature selection**)
- We will focus on **combining features**

Data Understanding

Name	Solubility	No. C atoms	Fraction of rotatable bonds	Topol. diam.	Geom. diam.	LogP	No. heavy bonds	...
methylpentane	good	6	0.40	4	3.46	2.44	5	...
methylcyclohexene	good	7	0	4	3.00	2.51	7	...
nonene	med.	9	0.75	8	6.93	3.53	8	...
hexadiene	good	6	0.60	5	4.36	2.14	5	...
butadiene	good	4	0.33	3	2.65	1.36	3	...
naphthalene	good	10	0	5	3.61	2.84	11	...
acenaphthylene	good	12	0	5	3.58	3.32	14	...
pyrene	poor	16	0	7	5.00	4.58	19	...
dimethylantracene	poor	16	0	7	5.29	4.61	18	...
hexahdropyrene	med.	16	0	7	5.00	3.82	19	...
triphenylene	poor	18	0	7	5.00	5.15	21	...
benzo(e)pyrene	poor	20	0	7	5.29	5.64	24	...

Plot the Data: two variables



Dimensionality Reduction

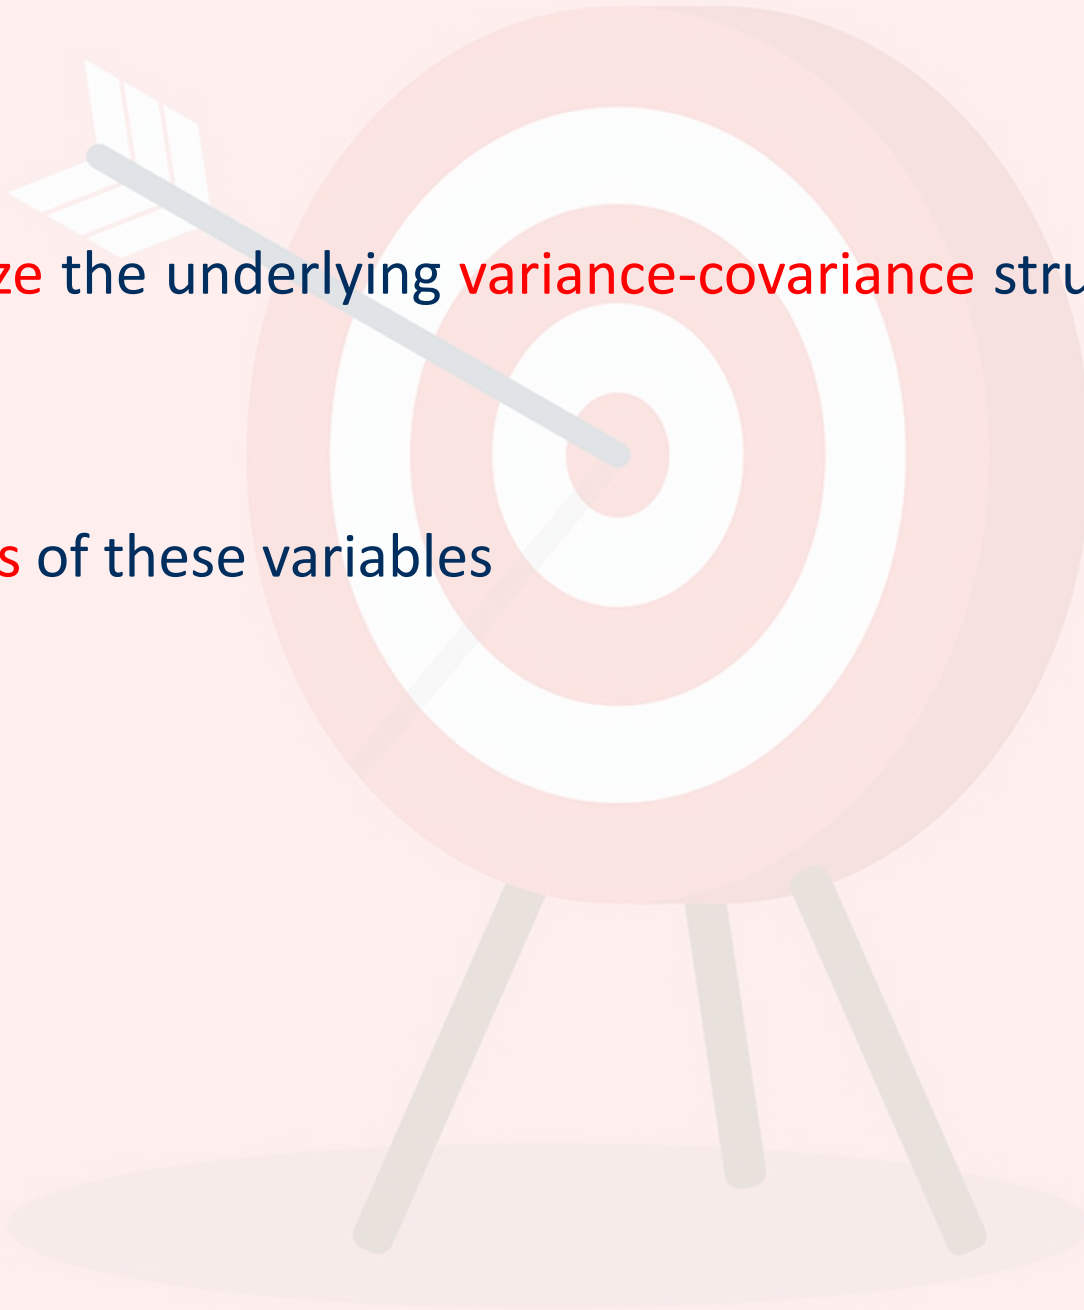
- Statistical methods that provide information about **point scatters** in multivariate space
- Simplify **complex relationships** between cases and/or variables
- Make it easier to **identify patterns**
- Remove **correlation** between features

Why reduce the dimensionality?

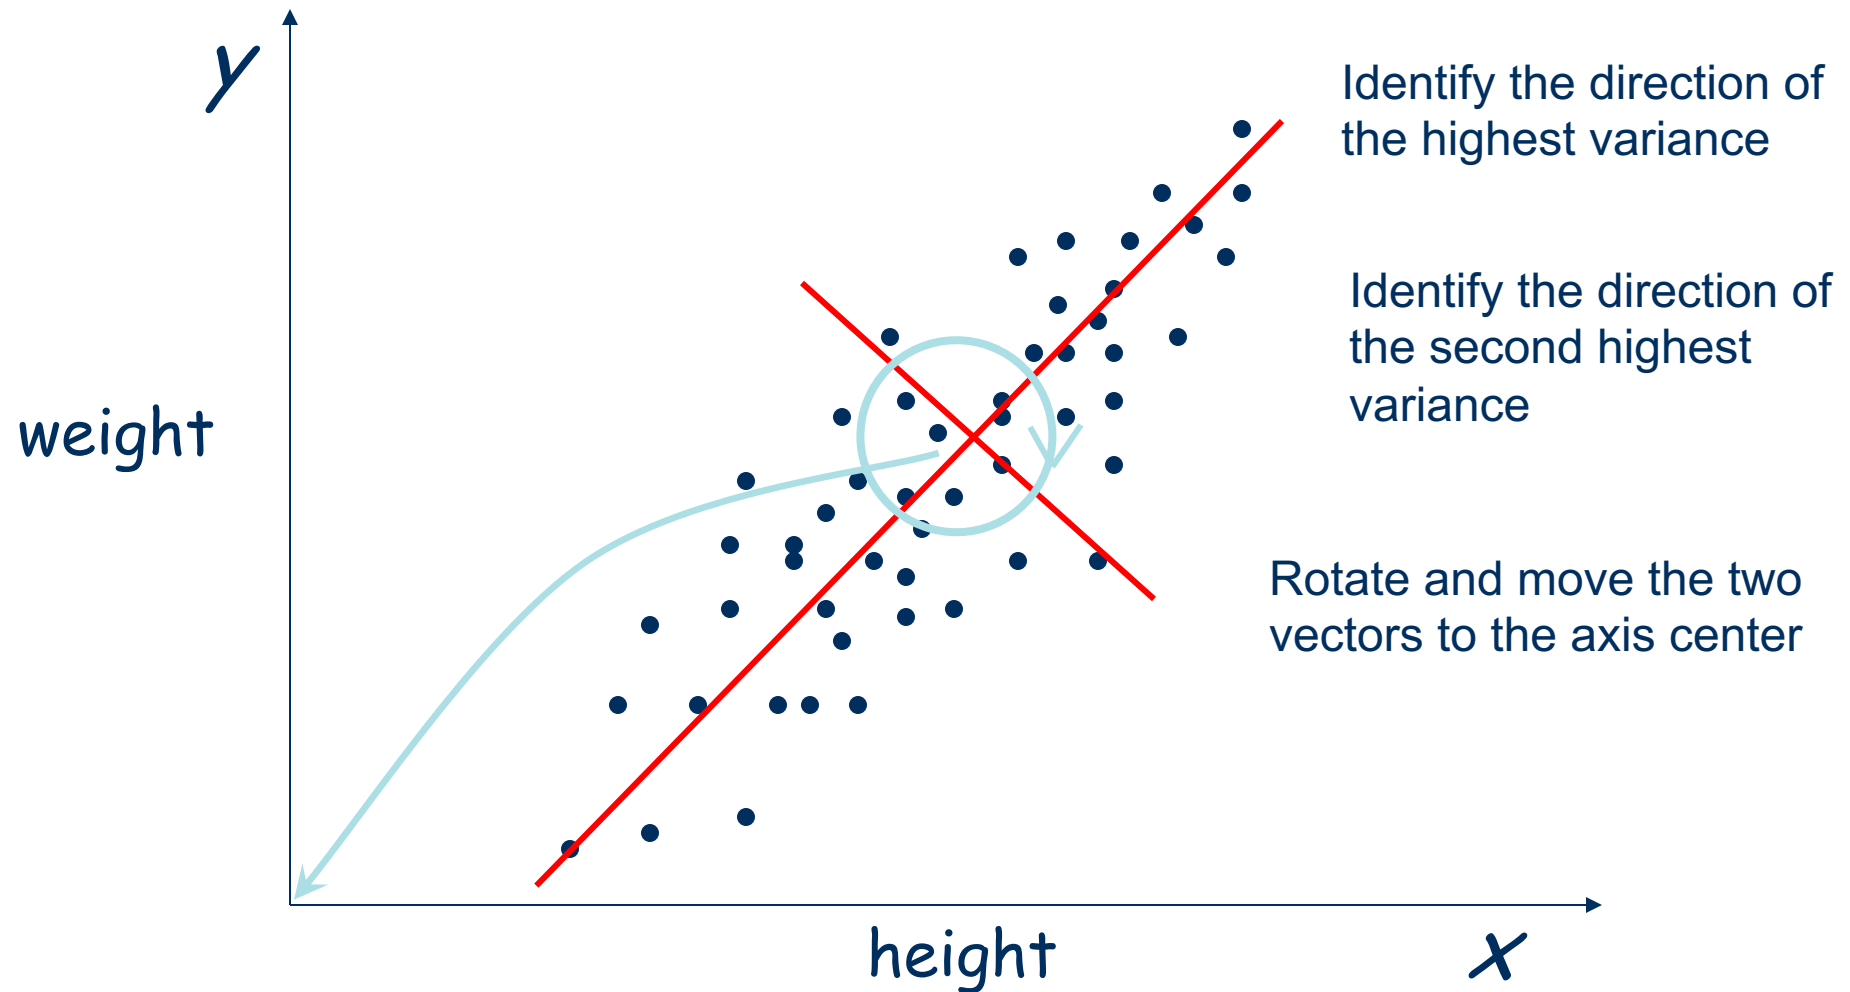
- Better presentation than ordinal axes
- Do we need a 100-dimensional space to view the data?
- Question: How to find the “best” low dimensional space that conveys maximum useful information?
- One answer: Find “Principal Components”

The goal


- We wish to **explain/summarize** the underlying **variance-covariance** structure of a large set of variables
- Use a few **linear combinations** of these variables



Imagine a two-dimensional scatter of points that show a high degree of correlation ...




Imagine a two-dimensional scatter of points that show a high degree of correlation ...



- Express the points using a **new coordinate system**
- The new axis corresponds to a linear combination of the dimensions of the original objects
- **OBS:** the points are spread mostly on the new x direction (Var 1), while the spread on direction y (Var 2) is low!


What does this mean?

- More “efficient” description
 - 1st Var captures max. variance
 - 2nd Var captures the max. amount of residual variance, at right angles (orthogonal) to the first
- The 1st Var may capture so much of the information content in the original data set that we can ignore the remaining axis



Principal Components Analysis (PCA)

- **Why:**
 - clarify relationships among variables
 - clarify relationships among objects
- **When:**
 - significant correlations exist among variables
- **How:**
 - define new axes (components)
 - examine **correlation** between axes and variables
 - find **scores** of objects on new axes



Philosophy of PCA

- We typically have a data matrix of M observations on N correlated variables

$$\mathbf{X} = \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$$

- PCA looks for a transformation of \mathbf{x}_i into N new variables $\mathbf{y} = \mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_N$ that are uncorrelated
- Dimensionality reduction implies **information loss**
- PCA preserves as much information as possible, that is, it **minimizes** the error

$$||\mathbf{x} - \mathbf{y}||$$

PCA: output

- New variables b_i that are linear combination of the original variables (x_i):

$$b_i = u_{i1}x_1 + u_{i2}x_2 + \dots + u_{iN}x_N, i=1\dots N$$

- The new variables b_i are derived in decreasing order of importance
- They are called “Principal Components”

PCA: more formally

From N original variables: x_1, x_2, \dots, x_N :

Produce N new variables: b_1, b_2, \dots, b_N :

$$b_1 = u_{11}x_1 + u_{12}x_2 + \dots + u_{1N}x_N$$

$$b_2 = u_{21}x_1 + u_{22}x_2 + \dots + u_{2N}x_N$$

...

$$b_N = u_{N1}x_1 + u_{N2}x_2 + \dots + u_{NN}x_N$$

such that:

u_N 's are uncorrelated (orthogonal)

u_1 explains as much as possible of the original variance in the data set

u_2 explains as much as possible of the remaining variance

...



PCA: Covariance Matrix

- **Question:** How should we determine the “best” lower dimensional space?
- **Answer:**

The “best” low-dimensional space can be determined by the “best” **eigenvectors** of the **covariance** matrix of the data (i.e., the eigenvectors corresponding to the “largest” **eigenvalues** – also called “Principal components”)

PCA: Covariance Matrix

$$\mathbf{X} = \begin{bmatrix} X_1 \\ \vdots \\ X_n \end{bmatrix}$$

$$\text{cov}(X_i, X_j) = E[(X_i - \mu_i)(X_j - \mu_j)]$$

$$\Sigma = \begin{bmatrix} E[(X_1 - \mu_1)(X_1 - \mu_1)] & E[(X_1 - \mu_1)(X_2 - \mu_2)] & \cdots & E[(X_1 - \mu_1)(X_n - \mu_n)] \\ E[(X_2 - \mu_2)(X_1 - \mu_1)] & E[(X_2 - \mu_2)(X_2 - \mu_2)] & \cdots & E[(X_2 - \mu_2)(X_n - \mu_n)] \\ \vdots & \vdots & \ddots & \vdots \\ E[(X_n - \mu_n)(X_1 - \mu_1)] & E[(X_n - \mu_n)(X_2 - \mu_2)] & \cdots & E[(X_n - \mu_n)(X_n - \mu_n)] \end{bmatrix}.$$

$$\text{cov}(X, Y) = \frac{1}{n} \sum_{i=1}^n (x_i - E(X))(y_i - E(Y))$$

Eigenvectors

$\{u_{11}, u_{12}, \dots, u_{1N}\}$: 1st **eigenvector** of the covariance matrix, and **coefficients** of the 1st principal component

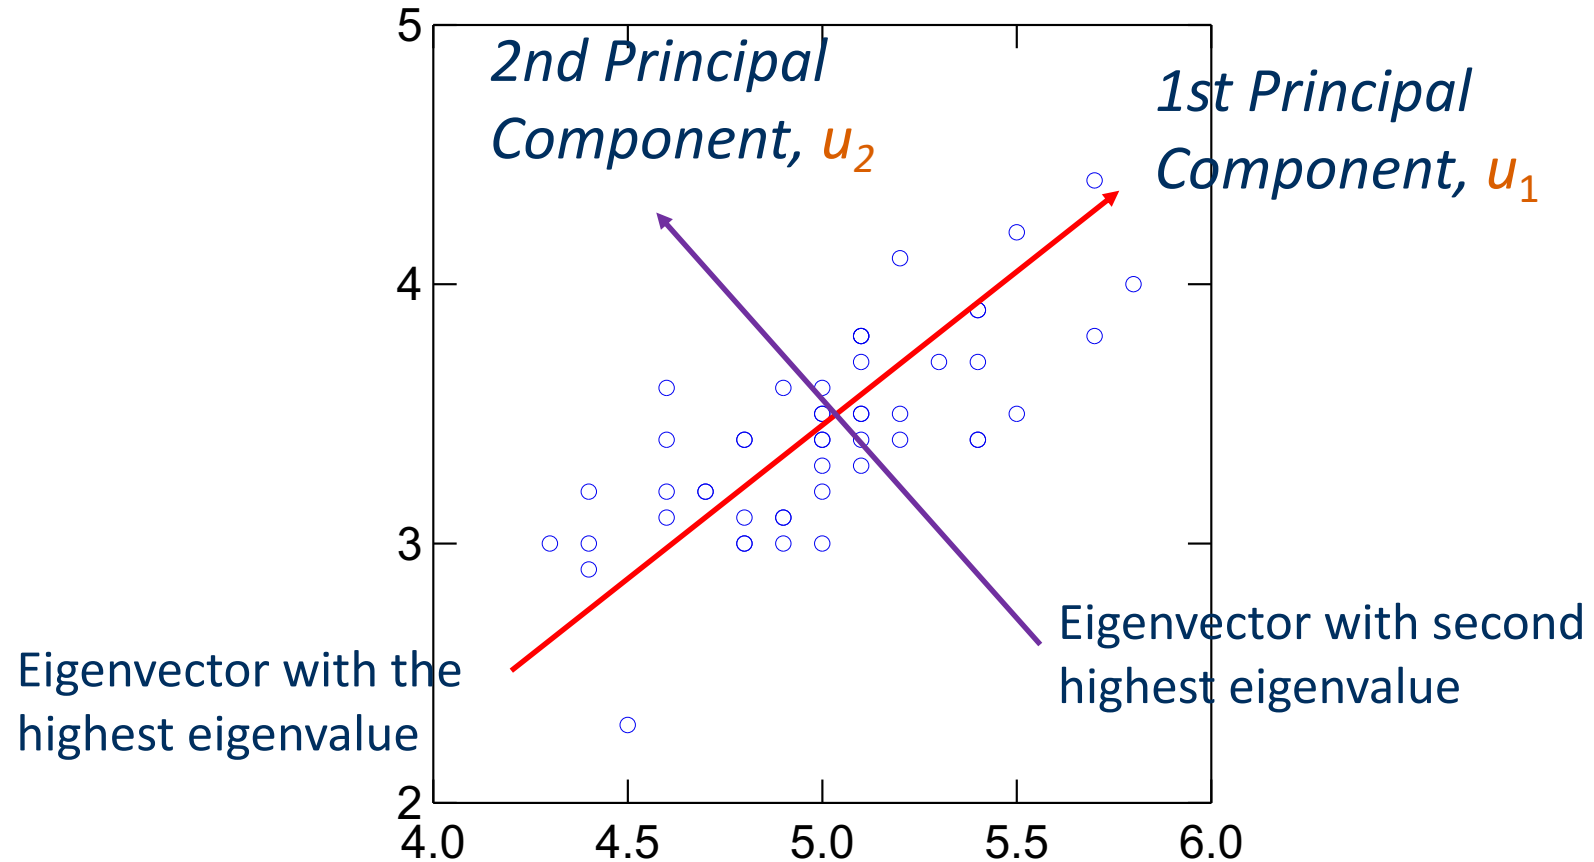
$\{u_{21}, u_{22}, \dots, u_{2N}\}$: 2nd **eigenvector** of the covariance matrix, and **coefficients** of the 2nd principal component

...


$\{u_{N1}, u_{N2}, \dots, u_{NN}\}$: N^{th} **eigenvector** of the covariance matrix, and **coefficients** of the N^{th} principal component




Singular Value Decomposition (SVD)



PCA Eigenvalues



PCA Scores



PCA Steps

Step 1: $\bar{x} = \frac{1}{M} \sum_{i=1}^M x_i$

Step 2: subtract the mean: $\Phi_i = x_i - \bar{x}$ (i.e., center at zero)

Step 3: form the matrix $A = [\Phi_1 \ \Phi_2 \ \cdots \ \Phi_M]$ ($N \times M$ matrix), then compute:

$$C = \frac{1}{M} \sum_{n=1}^M \Phi_n \Phi_n^T = AA^T$$

(sample **covariance** matrix, $N \times N$, characterizes the *scatter* of the data)

Step 4: compute the eigenvalues of C : $\lambda_1 > \lambda_2 > \cdots > \lambda_N$

Step 5: compute the eigenvectors of C : u_1, u_2, \dots, u_N



PCA Steps (cont'd)

- Since C is symmetric, u_1, u_2, \dots, u_N form a basis, (i.e., any vector x or actually $(x - \bar{x})$, can be written as a linear combination of the eigenvectors):

$$x - \bar{x} = b_1 u_1 + b_2 u_2 + \dots + b_N u_N = \sum_{i=1}^N b_i u_i \quad , \text{ where } b_i = \frac{(x - \bar{x}) \cdot u_i}{(u_i \cdot u_i)}$$

Step 6: (**dimensionality reduction step**) keep only the terms corresponding to the K largest eigenvalues:

$$\hat{x} - \bar{x} = \sum_{i=1}^K b_i u_i \text{ where } K \ll N$$

- The representation of $\hat{x} - \bar{x}$ into the basis u_1, u_2, \dots, u_K is thus

$$\begin{bmatrix} b_1 \\ b_2 \\ \dots \\ b_K \end{bmatrix}$$

PCA: Linear Transformation

- Every object \mathbf{x} in the original space is mapped as follows:

$$\begin{bmatrix} b_1 \\ b_2 \\ \dots \\ b_K \end{bmatrix} = \begin{bmatrix} u_1^T \\ u_2^T \\ \dots \\ u_K^T \end{bmatrix} (x - \bar{x}) = U^T (x - \bar{x})$$

Step 4: Singular Value Decomposition (SVD)*

$$A = \underset{[M \times r]}{U} \underset{[r \times r]}{S} \underset{[r \times N]}{V^T} = \begin{bmatrix} \vec{u}_1 & \vec{u}_2 & \cdots & \vec{u}_r \end{bmatrix} \begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_r \end{bmatrix} \begin{bmatrix} \vec{v}_1 \\ \vec{v}_2 \\ \vdots \\ \vec{v}_r \end{bmatrix}$$

- **r** : rank of matrix **A**
- $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_r$: singular values (square roots of eig-values AA^T , A^TA)
- $\vec{u}_1, \vec{u}_2, \dots, \vec{u}_r$: left singular vectors (eig-vectors of AA^T)
- $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_r$: right singular vectors (eig-vectors of A^TA)

$$A = \lambda_1 \vec{u}_1 \vec{v}_1^T + \lambda_2 \vec{u}_2 \vec{v}_2^T + \cdots + \lambda_r \vec{u}_r \vec{v}_r^T$$

Example SVD*

- Initial **matrix**:

$$A = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix}$$

- Transpose:

$$A^T = \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix}$$

- Then compute:

$$AA^T = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix} \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix}$$

- Find the **eigenvectors** and **eigenvalues**!

* Find the step-by-step solution here: <https://datajobs.com/data-science-repo/SVD-Tutorial-%5BKirk-Baker%5D.pdf>

Example SVD

- Matrix:

$$AA^T = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix} \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix}$$

- Finding eigenvectors and eigenvalues:

$$\begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \lambda \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- Solve the above system of equations!

Example SVD

- System of equations:

$$\begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \lambda \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- After solving it, we get:

$$\begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

- After normalizing:

$$U = \begin{bmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} & \frac{-1}{\sqrt{2}} \end{bmatrix}$$

Example SVD

- In a similar way we get **V**:

$$V^T = \begin{bmatrix} \frac{1}{\sqrt{6}} & \frac{2}{\sqrt{6}} & \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{5}} & \frac{-1}{\sqrt{5}} & 0 \\ \frac{1}{\sqrt{30}} & \frac{2}{\sqrt{30}} & \frac{-5}{\sqrt{30}} \end{bmatrix}$$

- We have also recorded the **eigenvalues of U and V** (they are the same!)

$$S = \begin{bmatrix} \sqrt{12} & 0 & 0 \\ 0 & \sqrt{10} & 0 \end{bmatrix}$$

Example SVD

- So finally, we have:

$$A_{mn} = U_{mm} S_{mn} V_{nn}^T = \begin{bmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} & \frac{-1}{\sqrt{2}} \end{bmatrix} \begin{bmatrix} \sqrt{12} & 0 & 0 \\ 0 & \sqrt{10} & 0 \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{6}} & \frac{2}{\sqrt{6}} & \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{5}} & \frac{-1}{\sqrt{5}} & 0 \\ \frac{1}{\sqrt{30}} & \frac{2}{\sqrt{30}} & \frac{-5}{\sqrt{30}} \end{bmatrix} =$$
$$\begin{bmatrix} \frac{\sqrt{12}}{\sqrt{2}} & \frac{\sqrt{10}}{\sqrt{2}} & 0 \\ \frac{\sqrt{12}}{\sqrt{2}} & \frac{-\sqrt{10}}{\sqrt{2}} & 0 \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{6}} & \frac{2}{\sqrt{6}} & \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{5}} & \frac{-1}{\sqrt{5}} & 0 \\ \frac{1}{\sqrt{30}} & \frac{2}{\sqrt{30}} & \frac{-5}{\sqrt{30}} \end{bmatrix} = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix}$$

More about SVD:

<https://datajobs.com/data-science-repo/SVD-Tutorial-%5BKirk-Baker%5D.pdf>

How to choose K?

- Choose **K** using the following criterion:

$$\frac{\sum_{i=1}^K \lambda_i}{\sum_{i=1}^N \lambda_i} > \textit{Threshold} \quad (\text{e.g., } 0.9 \text{ or } 0.95)$$

- In this case, we say that we “**preserve**” 90% or 95% of the information in the data
- If **K=N**, then we “preserve” 100% of the information in the data

Normalization

- The principal components are dependent on the *units* used to measure the original variables as well as on the *range* of values they assume
- Data should always be *normalized* prior to using PCA
- A common normalization method is to transform all the data to have *zero mean* and *unit standard deviation*:

$$\frac{x_i - \mu}{\sigma}$$

(μ and σ are the mean and standard deviation of x_i 's)


PCA Summary

- Rotates a multivariate dataset into a new configuration which is easier to interpret
- Purpose:
 - simplify data
 - look at relationships between variables
 - identify patterns in the correlated variables

Example PCA

Data =

2.5	2.4
0.5	0.7
2.2	2.9
1.9	2.2
3.1	3.0
2.3	2.7
2	1.6
1	1.1
1.5	1.6



Example PCA

- Subtract the mean (can also divide by standard deviation)

	x	y
	2.5	2.4
	0.5	0.7
	2.2	2.9
	1.9	2.2
Data =	3.1	3.0
	2.3	2.7
	2	1.6
	1	1.1
	1.5	1.6
	1.1	0.9

	x	y
	.69	.49
	-1.31	-1.21
	.39	.99
	.09	.29
DataAdjust =	1.29	1.09
	.49	.79
	.19	-.31
	-.81	-.81
	-.31	-.31
	-.71	-1.01

Example PCA

- Compute the **covariance matrix**:

$$cov = \begin{pmatrix} .616555556 & .615444444 \\ .615444444 & .716555556 \end{pmatrix}$$

- Compute the **eigenvalues** and **eigenvectors**:


$$eigenvalues = \begin{pmatrix} .0490833989 \\ 1.28402771 \end{pmatrix}$$

$$eigenvectors = \begin{pmatrix} -.735178656 & -.677873399 \\ .677873399 & -.735178656 \end{pmatrix}$$

Example PCA

- Plot

$$\begin{pmatrix} -.735178656 & -.677873399 \\ .677873399 & -.735178656 \end{pmatrix}$$



Example PCA

- Choose **eigenvectors** from:

$$FeatureVector = (eig_1 \ eig_2 \ eig_3 \ \ eig_n)$$

- Choose **both eigenvectors**:

$$\begin{pmatrix} -.677873399 & -.735178656 \\ -.735178656 & .677873399 \end{pmatrix}$$

- Or just the **first**:

$$\begin{pmatrix} -.677873399 \\ -.735178656 \end{pmatrix}$$


Example PCA

- Derive the **final data**:

$$FinalData = RowFeatureVector \times RowDataAdjust$$

- Hence, using **both eigenvectors**:

	x	y
Transformed Data=	-.827970186	-.175115307
	1.77758033	.142857227
	-.992197494	.384374989
	-.274210416	.130417207
	-1.67580142	-.209498461
	-.912949103	.175282444
	.0991094375	-.349824698
	1.14457216	.0464172582
	.438046137	.0177646297
	1.22382056	-.162675287



Example PCA


- Derive the **final data**:

$$FinalData = RowFeatureVector \times RowDataAdjust$$

- And using **one eigenvector**:

Transformed Data (Single eigenvector)

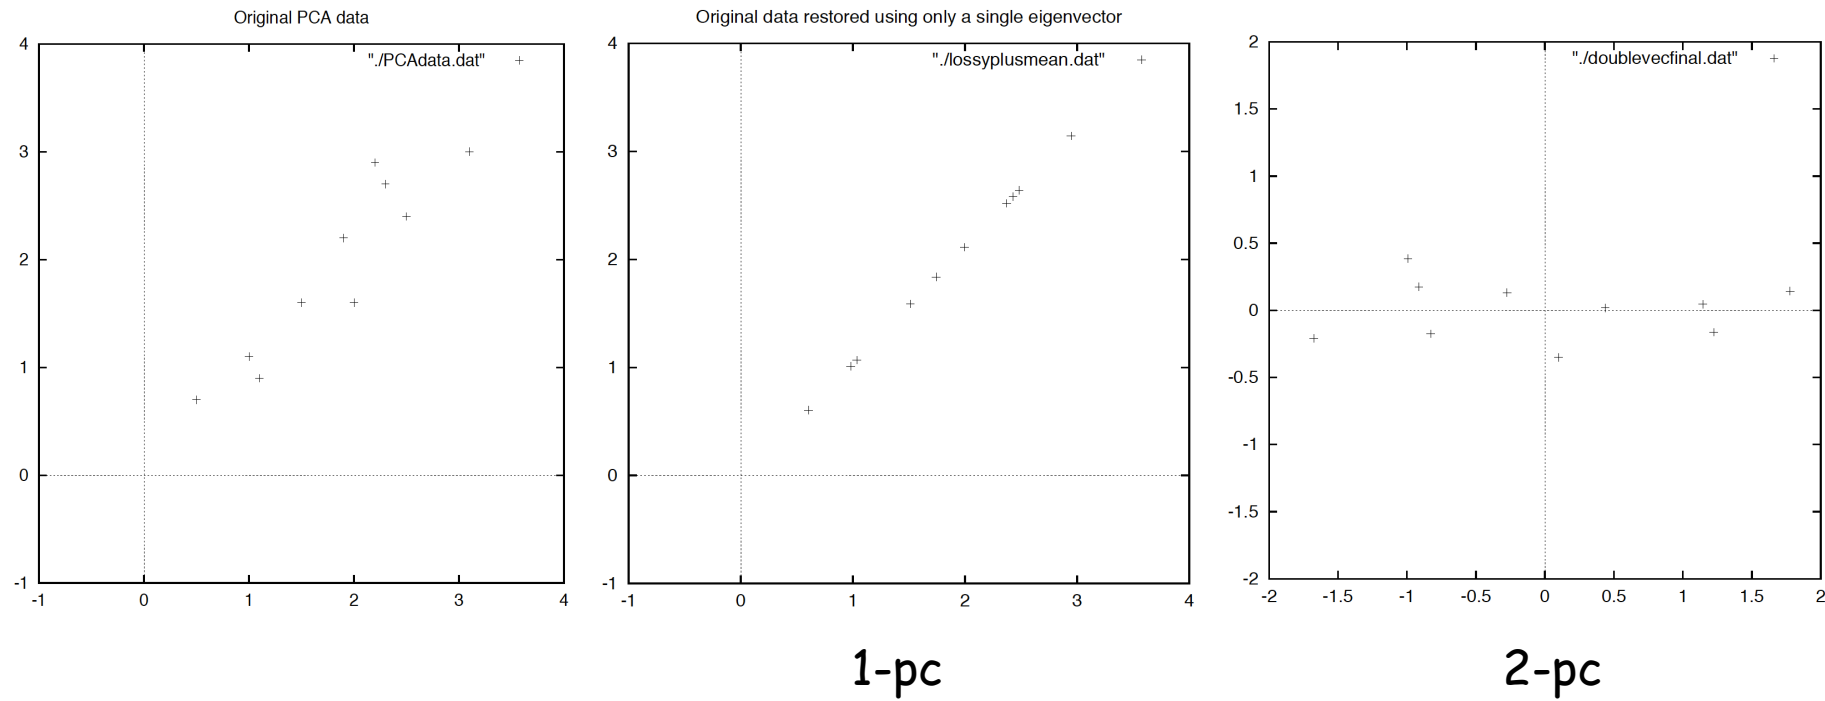
x
-.827970186
1.77758033
-.992197494
-.274210416
-1.67580142
-.912949103
.0991094375
1.14457216
.438046137
1.22382056




Example PCA

- Getting back the **original data**:

$$\text{RowOriginalData} = (\text{RowFeatureVector}^T \times \text{FinalData}) + \text{OriginalMean}$$



Multi-Dimensional Scaling (MDS)



- So far, we assumed that we know both data points **X** and distance matrix **D** between these points
- What if the original points **X** are not known but only distance matrix **D** is known?
- Can we reconstruct **X** or some approximation of **X**?

Problem


- Given distance matrix D between n points
- Find a k -dimensional representation of every x_i point i
- So that $d(x_i, x_j)$ is as close as possible to $D(i, j)$

Why do we want to do that?

Distances between 10 cities

	London	Berlin	Oslo	Moscow	Paris	Rome	Beijing	Istanbul	Gibraltar	Reykjavik
London	–									
Berlin	570	–								
Oslo	710	520	–							
Moscow	1550	1000	1020	–						
Paris	210	540	830	1540	–					
Rome	890	730	1240	1470	680	–				
Beijing	5050	4570	4360	3600	5100	5050	–			
Istanbul	1550	1080	1520	1090	1040	850	4380	–		
Gibraltar	1090	1450	1790	2410	960	1030	6010	1870	–	
Reykjavik	1170	1480	1080	2060	1380	2040	4900	2560	2050	–

Distances between 10 cities



Financial Indicators of Countries

Country	Increase	Life	IMR	TFR	GDP
Albania	1.2	69.2	30	2.9	659.91
Argentina	1.2	68.6	24	2.8	4343.04
Australia	1.1	74.7	7	1.9	17529.98
Austria	1.0	73.0	7	1.5	20561.88
Benin	3.2	45.9	86	7.1	398.21
Bolivia	2.4	57.7	75	4.8	812.19
Brazil	1.5	64.0	58	2.9	3219.22
Cambodia	2.8	50.1	116	5.3	97.39
China	1.1	66.7	44	2.0	341.31
Colombia	1.7	66.4	37	2.7	1246.87
Croatia	-1.5	67.1	9	1.7	5400.66
El Salvador	2.2	63.9	46	4.0	988.58
France	0.4	73.0	7	1.7	21076.77
Greece	0.6	75.0	10	1.4	6501.23
Guatemala	2.9	62.4	48	5.4	831.81
Iran	2.3	67.0	36	5.0	9129.34
Italy	-0.2	74.2	8	1.3	19204.92
Malawi	3.3	45.0	143	7.2	229.01
Netherlands	0.7	74.4	7	1.6	18961.90
Pakistan	3.1	60.6	91	6.2	385.59
Papua New Guinea	1.9	55.2	68	5.1	839.03
Peru	1.7	64.1	64	3.4	1674.15
Romania	-0.5	66.6	23	1.5	1647.97
USA	1.1	72.5	9	2.1	21965.08
Zimbabwe	4.4	52.4	67	5.0	686.75


Financial Indicators of Countries: First coordinate

Dim1: Measure of
overall development

Malawi	−2.027
Benin	−1.616
Cambodia	−1.414
Zimbabwe	−1.302
Pakistan	−1.133
Bolivia	−0.798
Papua New Guinea	−0.783
Guatemala	−0.706
El Salvador	−0.344
Peru	−0.277
Iran	−0.167
Brazil	−0.112
Colombia	0.036
China	0.188
Albania	0.220
Argentina	0.327
Romania	0.786
Greece	0.921
Australia	1.049
USA	1.105
Netherlands	1.158
Austria	1.164
Croatia	1.167
France	1.230

Financial Indicators of Countries: Two coordinates

Dim2: Measure GDP



How can we do that? (Algorithm)



High-level view of the MDS algorithm

- Randomly initialize the positions of n points in a k -dimensional space
- Compute pairwise distances D' for this placement
- Compare D' to D
- Move points to adjust their pairwise distances better (make D' closer to D)
- Repeat until D' is close to D

The MDS algorithm

- **Input:** $n \times n$ distance matrix **D**
- Random n points in the k -dimensional space (x_1, \dots, x_n)
- **stop = false**
- **while not stop**
 - **totalerror = 0.0**
 - For every pair of points i, j compute
 - $D'(i, j) = d(x_i, x_j)$
 - $\text{error} = (D(i, j) - D'(i, j)) / D(i, j)$
 - **totalerror += error**
 - $x_i = (x_i - x_j) / D'(i, j) * \text{error}$
 - **If totalerror small enough, stop = true**

Questions about MDS

- Running time of the MDS algorithm
 - $O(n^2I)$, where I is the number of iterations of the algorithm
- MDS does not guarantee that the **metric property** is maintained in D'
- Faster? Guarantee of metric property?

Today...


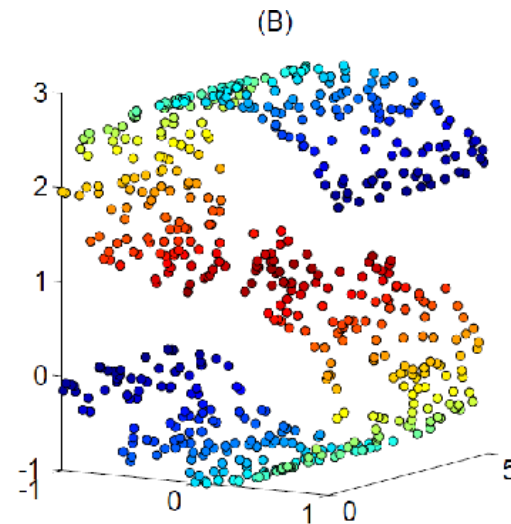

What is the
**Curse of
Dimensionality?**

The importance
of efficient **Data
Representation**

How do we
reduce Data
Dimensionality?

What is the **PCA**
algorithm?

What is the
MDS algorithm?




TODOs



Reading:

Main course book: Chapter 6
(Sec. 6.1, 6.3, 6.5)



Lab 1

Recommended to complete the lab
before the end of the week



Quiz 1

Coming up next

