

# Lecture 7

## Classification II

KNN, SVM, Neural networks

**Golnaz Taheri, PhD**

Senior Lecturer, Stockholm University



# Outline

---

- Eager vs Lazy learning
- The nearest neighbor classifier
- The Perceptron
- Support Vector Machine (SVM)
- Logistic regression

# Classification recap

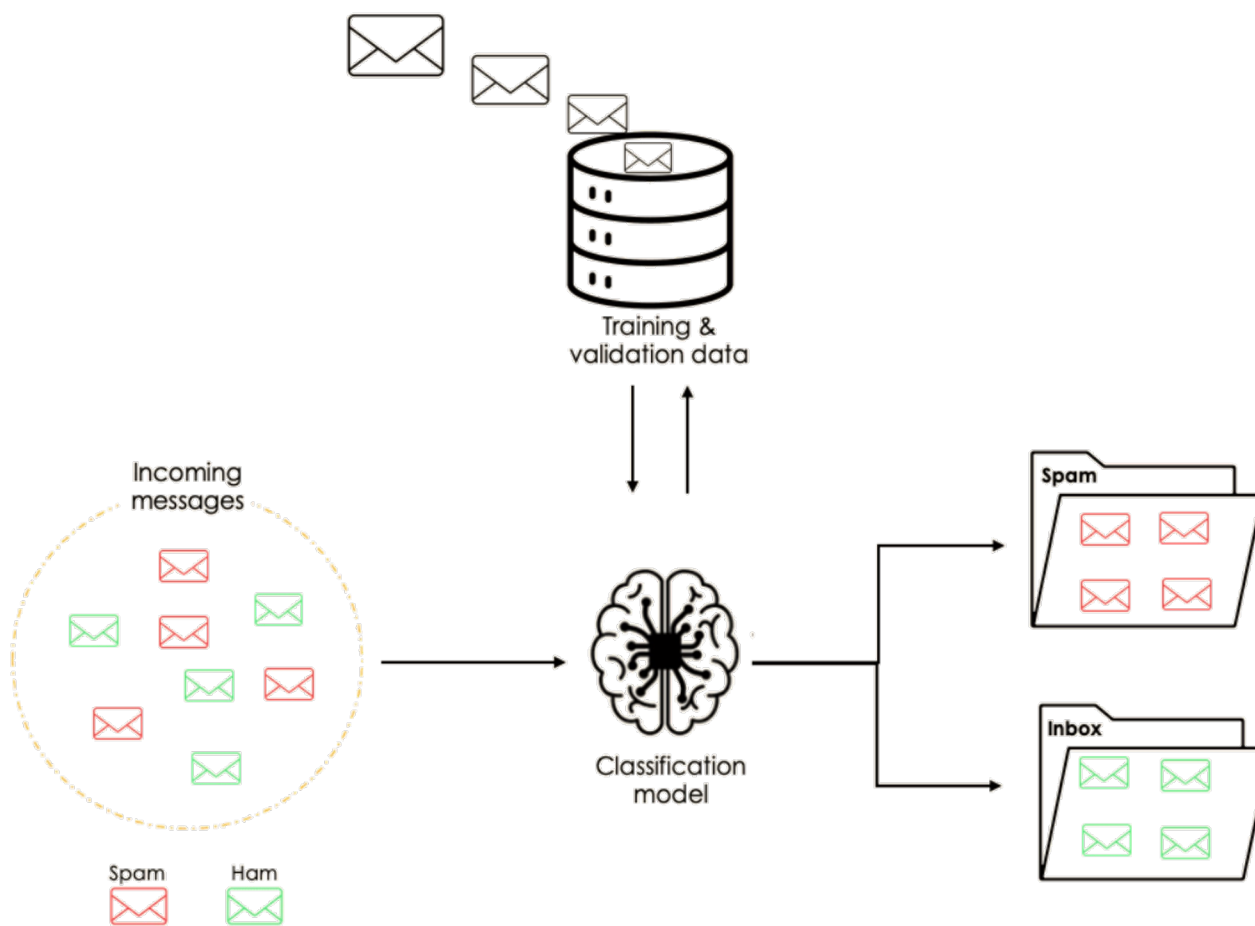
---

- Classification is a **supervised** machine learning method where the model tries to **predict** the correct **label** of a given input data.
- In classification, the model is **fully trained** using the training data, and then it is **evaluated** on **test data** before being used to perform prediction on new unseen data.



# Classification recap

- An algorithm can learn to predict whether a given email is spam or no spam.



# Eager Learning

---

- **Eager learning** methods construct **general, explicit** description of the target function based on the provided training examples.

≡ *one-fits-all*

≡ *input independent*

- Learn the **model** as soon as the **training data** becomes available
- More training time, less prediction time.
  - Support vector machines (SVM)
  - Decision tree
  - Artificial neural networks



# Lazy Learning

---

- **Lazy learning** methods simply store the data and generalizing beyond these data is postponed until an explicit request is made.
- Delay **model-building** until **testing data** needs to be classified
- Less training time, more prediction time.
- The model itself **consists of the training data**



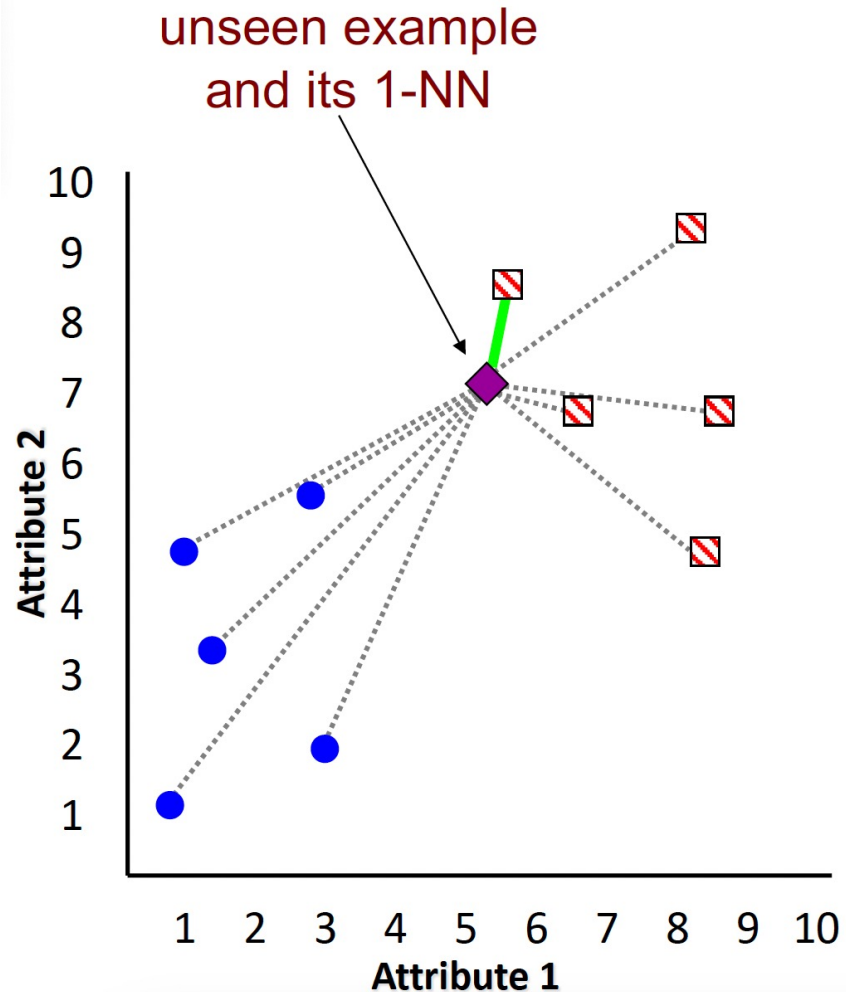
# K-Nearest Neighbor Classification

---

- **KNN** algorithm is one of the simplest classification algorithm
- **Non-parametric**
  - it does not make any assumptions on the underlying data distribution
- **Lazy** learning algorithm.
  - Means that the **training** phase is pretty **fast** .
  - **Lack** of **generalization** means that KNN keeps all the training data.
- Its purpose is to use a database in which the data points are separated into several classes to predict the classification of a new sample point.



# Nearest Neighbor Classifier




Goal: build **no model**, just find the **most similar** training example(s)





# The decision boundary of 1-NN


- **Decision boundaries:** surfaces that simply divide the space into regions “belonging “ to each Instance
- For 1-NN it is the **Voronoi Diagram** of the training set




# 1-NN is sensitive to outliers

## Solution: K-NN classifier:

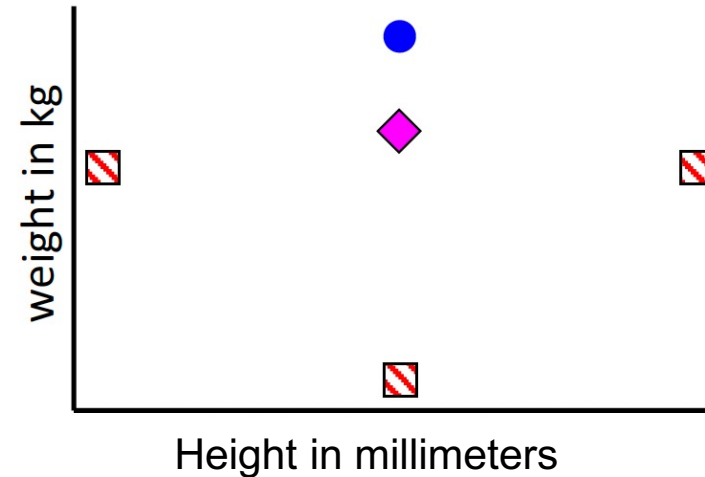
- We measure the distance to the **nearest K instances**, and let them **vote**
- K is typically chosen to be an **odd number**



# 1-NN is sensitive to the units of measurement



X axis measured in **centimeters**  
Y axis measure in **kg**  
The nearest neighbor to the **pink** unknown instance is **red**




X axis measured in **millimeters**  
Y axis measure in **kg**  
The nearest neighbor to the **pink** unknown instance is **blue**

**Solution:** normalize the units per attribute


- min-max normalization
- standardization



# 1-NN is sensitive to irrelevant attributes



1 attribute  
Can provide perfect  
classification



2 attributes  
Classification is  
wrong!

**Solution:** add more training data or ask domain experts  
what features are important




# Characteristics of K-NN classifiers


---

- **Lazy learners:** computational time in classification, no model building
- **Eager learners:** computational time in model building
- **Decision trees** try to find **global** models, K-NN classifiers take into account **local** information, i.e., example-based information, by looking at the neighborhood of the test example
- K-NN classifiers depend a lot on the choice of **distance measure**


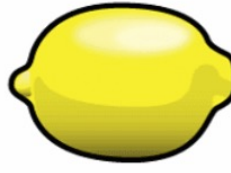




# Distance measure importance

iknowit  Classifying Objects by Color

Assign 

Which object belongs in this group?


  


Hint Submit

Progress 2/15

Score 2



Color
Red
Green
Blue
Yellow
Purple




Color
1
2
3
4
5





# Distance measure importance

iknowit  Classifying Objects by Color Assign 

Which object belongs in this group?


 

Progress 2/15  
Score 2



Hint Submit

Color
Red
Green
Blue
Yellow
Purple



Red	Green	Blue	Yellow	Purple
1	0	0	0	0
0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	1

# The Perceptron

---

- **Input:** each example  $x$  has a set of attributes  $x = \{a_1, a_2, \dots, a_m\}$  and is of class  $y$
- **Estimated classification output:**  $u$
- **Task:** express each sample  $x$  as a weighted combination (linear combination) of the attributes






# The Perceptron

- Define a set of  $m$  weights:  $\langle w_1, w_2, \dots, w_m \rangle$
- Multiply each attribute with their weight and sum them up
- Use an additional weight  $w_0$ 
  - intercept value, makes the model more general

$$f(x) = w_1 a_1 + w_2 a_2 + \dots + w_m a_m + w_0 = \underbrace{\langle w, x \rangle}_{\text{inner or dot product of } w \text{ and } x} + w_0$$



$$f(x) > 0 \rightarrow u_i = 1$$

$$f(x) \leq 0 \rightarrow u_i = -1$$

How do you find the weights ?

# Online Perceptron Algorithm

---

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$

Repeat

Accept training example  $i : (\mathbf{x}_i, y_i)$

$$u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$$

if  $y_i \cdot u_i \leq 0$

$$\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$$



# The Perceptron: example

---

4 examples with 4 attributes

- Suppose our training set contains 4 examples with 3 attributes:

$\alpha_1$	$\alpha_2$	$\alpha_3$	Class Y
0	0	0	1
0	1	0	1
1	0	1	-1
0	1	1	-1

- We want to learn a straight line (hyperplane in this case) that separates them:

$$f(x) = \langle w, x \rangle + w_0 = w_1\alpha_1 + w_2\alpha_2 + w_3\alpha_3 + w_0$$



# The Perceptron: example

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$   
 Repeat  
     Accept training example  $i : (\mathbf{x}_i, y_i)$   
      $u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$   
     if  $y_i \cdot u_i \leq 0$   
          $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

W0	W1	W2	W3		$\alpha_0$	$\alpha_1$	$\alpha_2$	$\alpha_3$	y	f(x)	class
0	0	0	0	$\mathbf{x}_1$	1	0	0	0	1	0	$\rightarrow -1$
					1	0	1	0	1		
					1	1	0	1	-1		
					1	0	1	1	-1		

---

W0	W1	W2	W3
1	0	0	0

$$f(x) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0$$

$$f(x) \leq 0 \rightarrow u_i = -1$$

$$f(x) > 0 \rightarrow u_i = 1$$



# The Perceptron: example

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$   
 Repeat  
   Accept training example  $i : (\mathbf{x}_i, y_i)$   
    $u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$   
   if  $y_i \cdot u_i \leq 0$   
      $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

W0	W1	W2	W3	$\alpha_0$	$\alpha_1$	$\alpha_2$	$\alpha_3$	y	f(x)	class
1	0	0	0	1	0	0	0	1	1	$\rightarrow 1$
				1	0	1	0	1	1	$\rightarrow 1$
				$\mathbf{x}_3$ 1	1	0	1	-1	<b>1</b>	<b><math>\rightarrow 1</math></b>
				1	0	1	1	-1		

---

W0	W1	W2	W3
0	-1	0	-1

$$f(x) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0 = 0$$

$$f(x) \leq 0 \rightarrow u_i = -1 \qquad f(x) > 0 \rightarrow u_i = 1$$



# The Perceptron: example

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$   
Repeat  
  Accept training example  $i : (\mathbf{x}_i, y_i)$   
   $u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$   
  if  $y_i \cdot u_i \leq 0$   
     $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

w0	w1	w2	w3
1	0	0	0

	$\alpha_0$	$\alpha_1$	$\alpha_2$	$\alpha_3$	$y$	$f(x)$	class
	1	0	0	0	1	1	$\rightarrow 1$
	1	0	1	0	1	1	$\rightarrow 1$
$\mathbf{x}_3$	1	1	0	1	-1	1	$\rightarrow 1$
	1	0	1	1	-1		

w0	w1	w2	w3
0	-1	0	-1

$$f(x) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0 = 0$$

$$f(x) \leq 0 \rightarrow u_i = -1 \qquad f(x) > 0 \rightarrow u_i = 1$$



# The Perceptron: example

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$   
 Repeat  
     Accept training example  $i: (\mathbf{x}_i, y_i)$   
      $u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$   
     if  $y_i \cdot u_i \leq 0$   
          $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

W0	W1	W2	W3		$\alpha_0$	$\alpha_1$	$\alpha_2$	$\alpha_3$	y	f(x)	class
0	-1	0	-1	$\mathbf{x}_1$	1	0	0	0	1	0	$\rightarrow -1$
					1	0	1	0	1		
					1	1	0	1	-1		
					1	0	1	1	-1	-1	$\rightarrow -1$

---

W0	W1	W2	W3
1	-1	0	-1

$$f(x) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0 = 0$$

$$f(x) \leq 0 \rightarrow u_i = -1 \qquad f(x) > 0 \rightarrow u_i = 1$$



# The Perceptron: example

$f(x)$  can also be written as a linear combination of all or some of the **training examples**

Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$

Repeat

Accept training example  $i : (\mathbf{x}_i, y_i)$

$u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$

if  $y_i \cdot u_i \leq 0$

$\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

w0	w1	w2	w3	$\alpha_0$	$\alpha_1$	$\alpha_2$	$\alpha_3$	y	f(x)	class
1	-1	0	-1	1	0	0	0	1	1	$\rightarrow 1$
				1	0	1	0	1	1	$\rightarrow 1$
				1	1	0	1	-1	-1	$\rightarrow -1$
				1	0	1	1	-1	-0	$\rightarrow -1$

$$f(x) = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0 = -\alpha_1 - \alpha_3 + 1$$

$$f(x) = \sum_i k_i y_i \langle x_i, x \rangle + b$$

$$k_1 = 2 \quad k_2 = 0 \quad k_3 = 1 \quad k_4 = 0$$

$$f(x) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = w_1 \alpha_1 + w_2 \alpha_2 + w_3 \alpha_3 + w_0 = 0$$

$$f(x) \leq 0 \rightarrow u_i = -1$$

$$f(x) > 0 \rightarrow u_i = 1$$






# The Perceptron

---

$$w_0 + w_1 a_1 + w_2 a_2 + \dots + w_m a_m = 0$$




The perceptron learning algorithm is guaranteed to find a separating hyperplane – if there is one.



# Many possible separating hyperplanes


---



Which hyperplane to choose?


# Choosing hyperplanes

---



Margin: defined by the two points (one from the '+' set and the other from the '-' set) with the minimum distance to the separating hyperplane

# Linear SVM



$x_1$ ,  $x_2$ ,  $x_3$ , and  $x_4$  are the **support vectors**, i.e., the closest points to the margins from both sides

# The maximum margin hyperplane


---

- One reasonable choice for the separating hyperplane:
  - The hyperplane that represents the largest separation, or **margin**, between the two classes
  - We choose the hyperplane so that the distance from it to the nearest data point(s) on each side is maximized
  - If such a hyperplane exists, it is known as the **maximum-margin hyperplane**
  - This is a **quadratic programming problem** that can be easily solved analytically by standard methods (optimization using Lagrange multipliers)



# The maximum margin hyperplane

- It is desirable to design linear classifiers that **maximize the margins** of their decision boundaries
- This ensures that their **worst-case generalization errors** are **minimized**



**Risk of overfitting is reduced by finding the maximum margin hyperplane!**



# The maximum margin hyperplane

---

- **Linear support vector machines:**

- Find the maximum-margin hyperplane
- Express this hyperplane as a linear combination of the data points  $(x_i, y_i)$ :


$$f(x) = \sum_i k_i y_i \langle x_i, x \rangle + b$$

- The points (from each side) with the smallest distance from the hyperplane are called **support vectors**
- **Formally:** support vectors are the elements of the training set that would change the position of the dividing hyperplane if removed
- We want to express  $f(x)$  as a linear combination of the support vectors  $x_i$  by identifying them and learning their corresponding weights  $k_i$



# Overfitting and slack variables


- **Curved decision boundary:** it is not always very likely to find a line dividing the training examples
- Data points may be noise or outliers
- **Alternative:** introduce a smooth boundary that ignores a few data points rather than defining a very curved boundary
- This is handled by introducing **slack variables  $\xi_i$**  that measure the distance of a point to its marginal hyperplane



**Prevention of overfitting!**



# What if the classes are not linearly separable?




Define a *mapping*  $\phi(x)$  that maps each attribute of point  $x$  to a new space where points *are linearly separable*

# What if the classes are not linearly separable?

- Original attributes are mapped to a new space to obtain linear separability

$$f_{orig}(x) = \langle w, x \rangle + w_0 = \sum_i k_i y_i \langle x_i, x \rangle + b$$

$$f(x) = \sum_i k_i y_i \langle \varphi(x_i), \varphi(x) \rangle + b$$



- Another alternative: **Artificial Neural Networks – deep learning**



# Kernels

- Assume two points in the original space:

$$x_1 = (a_1, a_2) \text{ and } x_2 = (b_1, b_2)$$

## Complex mapping:

- $\varphi(x)$  can be quite complex, for example:  $\varphi(x) = (a_1^2, a_2^2, \sqrt{2}a_1a_2)$

- Hence:

$$\varphi(x_1) = \varphi((a_1, a_2)) = (a_1^2, a_2^2, \sqrt{2}a_1a_2) \quad \varphi(x_2) = \varphi((b_1, b_2)) = (b_1^2, b_2^2, \sqrt{2}b_1b_2)$$

$$\langle \varphi(x_1), \varphi(x_2) \rangle = \langle (a_1^2, a_2^2, \sqrt{2}a_1a_2), (b_1^2, b_2^2, \sqrt{2}b_1b_2) \rangle \leftarrow \text{COMPLEX}$$

$$= (a_1^2b_1^2 + a_2^2b_2^2 + 2a_1b_1a_2b_2)$$

$$= (a_1b_1 + a_2b_2)^2$$

$$= \langle (a_1, a_2), (b_1, b_2) \rangle^2 = \langle x_1, x_2 \rangle^2$$

$\leftarrow$  SIMPLE

**kernel function:**  $K(x_1, x_2) = \langle \phi(x_1), \phi(x_2) \rangle$

A function that returns the value of the dot product between the images  $\varphi$  of two data points



# Kernels

## Observation:

- $\varphi(x)$  can be quite complex, for example see below

## kernel function:

$$K(x_1, x_2) = \langle \phi(x_1), \phi(x_2) \rangle$$

## Kernel trick:

- Do not need to actually compute the vectors  $\varphi$  in the mapped space
- Just need to identify a simple form of the dot product of the mapped values  $\varphi(x_1)$  and  $\varphi(x_2)$

$$\varphi(x_1) = \varphi((a_1, a_2)) = (a_1^2, a_2^2, \sqrt{2}a_1a_2) \quad \varphi(x_2) = \varphi((b_1, b_2)) = (b_1^2, b_2^2, \sqrt{2}b_1b_2)$$

$$\langle \varphi(x_1), \varphi(x_2) \rangle = \langle (a_1^2, a_2^2, \sqrt{2}a_1a_2), (b_1^2, b_2^2, \sqrt{2}b_1b_2) \rangle \leftarrow \text{COMPLEX}$$

$$= (a_1^2b_1^2 + a_2^2b_2^2 + 2a_1b_1a_2b_2)$$

$$= (a_1b_1 + a_2b_2)^2$$

$$= \langle (a_1, a_2), (b_1, b_2) \rangle^2 = \langle x_1, x_2 \rangle^2$$

$\leftarrow$  SIMPLE



# Beyond the Hyperplane, Choose your kernel


---

- **Linear (dot) kernel**

- This is linear classifier, use it as a test of non-linearity
- Or as a reference for the classification improvement with non-linear kernels

$$K(x_1, x_2) = \langle x_1, x_2 \rangle^1$$

**Linear Kernel**



Standard SVM with  
Hyperplane Boundary



# Beyond the Hyperplane, Choose your kernel


---

- **Polynomial**

- Simple, efficient for non-linear relationships
- Identifies/exploits polynomial relationships between the variables
- $d$  – degree, high  $d$  leads to overfitting

$$K(x_1, x_2) = \langle x_1, x_2 \rangle^d$$

**Polynomial Kernel**



More flexible, but more potential for overfitting



# Kernels

---

- **Several Types of Kernels:**

$$K(x_1, x_2) = \langle x_1, x_2 \rangle^d$$

polynomial kernels

$$K(x_1, x_2) = (\langle x_1, x_2 \rangle + 1)^d$$

$$K(x_1, x_2) = e^{-\|x_1 - x_2\|^2 / 2\sigma^2}$$


Gaussian radial basis function

$$K(x_1, x_2) = \tanh(\kappa \langle x_1, x_2 \rangle - \delta)$$

two-layer sigmoidal neural network

# Polynomial kernel: overfitting as $d$ increases!

---





# Which kernel?

---

- So **which kernel** and **which parameters** should I use?
- The answer is data-dependent
- Several kernels should be tried
- Try the **linear kernel** first
- Check if classification can be improved with **polynomial kernels**
- Then try other **nonlinear kernels** (tradeoff between quality of the kernel and the number of dimensions)
- **Select kernel + learn parameters**



# Logistic regression

---

- The logistic regression algorithm finds the best **logistic function** that can describe the relationship between two variables:
  - dependent variable **y** (class variable)
  - independent variable(s) **X** (data variables)
- **Classic logistic regression:**
  - **y** is **binary**: i.e., it has two possible outcomes, e.g., win/loss, health/unhealthy
  - since **y** is binary, we often label classes as either 1 or 0



# Computing the odds

---

- As **new examples** appear, we use the input variables and the logistic relationship to predict the **probability**  $p$  of a new example to belong to class  $y = 1$ :

$$P(y = 1 \mid X) = p$$

- Since  $p$  ranges between 0 and 1, we can convert this to a classification problem by using a **cutoff threshold**
- The higher the value of  $p$ , the more likely the new example belongs to class  $y = 1$ , instead of  $y = 0$
- For example, if we set the **cutoff** to 0.5, this would mean that an example will be classified as  $y = 1$ , when  $p > 0.5$ , otherwise it will be classified as  $y = 0$



# Odds ratio

---

- Odds refer to the **ratio** of the probability of an event to happen divided by the probability of not happening
- It is a metric representing the **likelihood of an event to occur**

$$\text{odds} = \frac{\text{probability of something happening}}{\text{probability of something not happening}}$$


- Hence, the odds of the observation belonging to class  $y = 1$  is  $p/(1 - p)$ :
  - When the odds are less than 1, they are against the example belonging to  $y = 1$
  - When the odds are greater than 1, they are for the example belonging to  $y = 1$



# Odds ratio

---

- For example, consider a fair six-sided dice:
  - The probability of a six coming up after a single roll is  $1/6$ , and  $5/6$  of not happening
  - The odds in favor of winning are  $(1/6) / (5/6) = 1/5$  or 1:5
  - The odds of losing are  $(5/6) / (1/6) = 5:1$
  - The odds are clearly against winning



# The logit function

---

- The **log odds function** or **logit function** is  $\log(p/(1 - p))$ , i.e., it corresponds to the logarithm of odds (the natural logarithm is most often used)
- **Ranges:**
  - $p$  ranges from 0 to 1
  - $p/(1 - p)$  ranges from 0 to *infinity*
  - $\log(p/(1 - p))$  ranges from *-infinity to infinity*

# Logistic regression

---

- Assume a set of ***m*** variables  $\{\alpha_1, \alpha_2, \dots, \alpha_m\}$  and a set of ***m*** weights  $\{w_1, w_2, \dots, w_m\}$
- Similar to the ***Perceptron*** formulation, we want to express the logit function as a **linear combination** of these variables:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$$

- The **logistic function** is the inverse of the logit function above, i.e., we solve for  $p$ :

$$p = \frac{1}{1 + e^{-(w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m)}}$$



# The sigmoid function


- The **logistic function** obtained above:

$$p = \frac{1}{1 + e^{-(w_0 + w_1 a_1 + w_2 a_2 + \dots + w_m a_m)}}$$

is a type of a **sigmoid function**:

$$\text{sigm}(h) = \frac{1}{1 + e^{-h}}$$

$$h = w_0 + w_1 a_1 + w_2 a_2 + \dots + w_m a_m$$



- The function ranges between 0 and 1



# Solving the logistic equation

---


- The goal is to find the values for the weights  $\{w_0, w_1, w_2, \dots, w_m\}$
- **Solution:** Maximum Likelihood Estimation (MLE)
- Use a likelihood function that measures how well a set of parameters fit a sample of data
- The parameter values that maximize the likelihood function are the maximum likelihood estimates
- Thus, the goal is to make inferences about the population that is most likely to have generated the training dataset
- **Assumption:** the data variables (attributes) are mutually independent



# Interpretation of logistic regression

- Suppose we have **one variable** and **five** examples

Input $x_1$	Binary Output $y$
0.5	0
1.0	0
0.65	0
0.75	1
1.2	1



- After solving the logistic equation, we get

$$\log(\text{odds}) = -4.411 + \mathbf{4.759}x_1$$

- This means that a one-unit increase of  $x_1$ , the log odds is expected to increase by **4.759**



# Interpretation of logistic regression

- Suppose we have **one variable** and **five** examples

Input $x_1$	Binary Output $y$
0.5	0
1.0	0
0.65	0
0.75	1
1.2	1

Given an input example with  $x_1 = 0.9$ , we get

$$p = \frac{1}{1 + e^{-(-4.411 + 4.758 * 0.9)}} = 46.8\%$$

If our **threshold is 0.5**, then the predicted class is  **$y = 0$**

- After solving the logistic equation, we get

$$\log(odds) = -4.411 + 4.759x_1$$

- This means that a one-unit increase of  $x_1$ , the log odds is expected to increase by **4.759**



# Types of logistic regression

---

- **Binary Logistic Regression:** the target variable has two possible categories; e.g., yes or no, spam or no spam, pass or fail
- **Multinomial Logistic Regression:** the target variable has three or more categories which are not in any particular order; e.g., categories of fruit: apple, mango, orange, and banana
- **Ordinal Logistic Regression:** the target variable has three or more ordinal categories; e.g., poor, average, good, very good, and excellent performance



# Assumptions of logistic regression


---

- The **target variable** should be binary, multinomial, or ordinal in nature
- The **observations should be independent** of each other; they should not come from repeated measurements
- There should be **little or no multicollinearity** among the independent variables; this means that the independent variables should not be too highly correlated with each other
- **Linearity** of the independent variables and log odds is assumed
- The success of Logistic Regression depends on the sample sizes; it requires a **large sample size** to achieve high predictive performance



# TODOs

---




## Reading:

Main course book chapters:

8.2.3, 8.6,

11.1-11.3,

14.1-14.5



## Lab 3

Sep 25



## Quiz 3



Stockholms  
universitet

# Coming up next

---


## Thursday

Lab 2 – Clustering using Python

Lecture 8 – Classification III

## Friday





Thanks!



`golnaz.taheri@dsv.su.se`

