


# Ensemble learning

Machine Learning (ML)

# Agenda

- Motivation
- Bias and Variance revisited
- Bagging
- Random Forests
- Boosting
- Stacking

# Motivation



= 1 inch (2,54 cm)

# Motivation

“If each member of a jury is more likely to be right than wrong, then the majority of the jury, too, is more likely to be right than wrong; and the probability that the right outcome is supported by a majority of the jury is a (swiftly) increasing function of the size of the jury, converging to 1 as the size of the jury tends to infinity.” - Marquis de Condorcet, 1785


# Bias and Variance revisited

- Given a **target** function  $f$
- **Find** a hypothesis  $h$ , such that
  - Such that  $h \approx f$
- **Bias** is the **tendency** of  $h$  to **deviate** from expected value when averaged on different training sets
  - High bias, linear model
  - Low bias, k-nearest neighbor
- **Variance** is the **amount** of change in  $h$  **due** to **fluctuations** in training data
  - Low variance, linear model (stable)
  - High variance, decision tree (unstable)
- **Bias-variance tradeoff**, the **choice** between a more **complex** low bias  $h$  vs **simple** low variance  $h$


# Ensemble learning

- **Successful** ensembles require **diversity**
  - Classifiers should make **different** mistakes
- 1. **Adjusting** induction method
  - By using **different** base **learners** (with different biases)
    - This will most **probably** ensure that they make **different errors**
- 2. Adjust **data** by using **homogenous** induction method
  - Bagging
  - Boosting
- 3. **Adjust** data **and** selection method
  - Random Forest

# Ensemble learning



versus



Traditional approach

- Build one **good** model

Ensemble approach

- Build **many** models and average the results

# Bagging

- Idea: Combining **many unstable** predictors to produce a **stable** predictor, by reducing **variance**
- A bootstrap sample  $B$  (one bag) of a set of examples  $E$  is created by **randomly selecting**  $n = |E|$  examples from  $E$  **with replacement**

- The **probability** of an example in  $E$  **appearing** in  $B$  is

$$1 - \left(1 - \frac{1}{n}\right)^n = 1 - \frac{1}{e} \approx 0.632$$

- Examples in  $E \setminus B$  is called **out-of-bag** examples
  - i.e. examples not in  $B$




## Bootstrap example

```
>>>from sklearn.utils import resample
>>>e = np.array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
>>>samp = resample(e)
>>>samp
array([1, 4, 8, 1, 5, 2, 8, 2, 3, 9])
>>>samp = resample(e)
>>>samp
array([3, 7, 7, 9, 2, 4, 4, 5, 2, 9])
>>>samp = resample(e)
>>>samp
array([1, 7, 2, 4, 8, 8, 8, 2, 1, 8])
>>>samp = resample(e)
>>>samp
array([9, 3, 9, 1, 6, 6, 9, 8, 6, 9])
>>>samp = resample(e)
>>>samp
array([6, 9, 6, 8, 5, 0, 0, 5, 9, 5])
```


# Bagging

- Take repeated **bootstrap samples** from training set  $D$
- Bootstrap sampling
  - Given training set  $D$  containing  $N$  training examples
    - Create  $D'$  by drawing  $N$  examples at random with replacement from  $D$
- Bagging\*
  - Create  $k$  bootstrap samples  $B_1, \dots, B_k$
  - Train a classifier (or regressor) on each  $B_i$
  - Predict new instance
    - by majority vote (classification)
    - by averaging value (regression)

\* L. Breiman. 1996. Bagging predictors. Machine Learning, 24(2):123-140



# Bagging example result



Breast-cancer-Wisconsin dataset from UCI repository  
Stratified 10 fold cross validation  
Decision trees as base classifier

## Bagging tips

- **Any** base learner can be used **but**,
  - good results are obtained by unstable (**high variance**) learners like
    - Decision trees
    - Neural networks
- Works in similar way for both classification and regression, only the **voting** scheme needs to change
- Using decision trees, **pruning** should **not** be used,
  - as this **decrease** the variance
- Can be used **without** separate **validation** or **test set**
  - **Out-of-bag** examples can be used for this

# Random Forest

- Idea: Combine **bagging** and **the random subspace method (RSM)** (feature bagging), in one method to further increase diversity (variance) in the base learner to produce a low **variance** predictor (ensemble).
- Similarly to bagging RSM samples features with replacement, typically the  $\sqrt{\text{no\_features}}$  is used
- $\text{No\_features} = 10, \sqrt{10} = 3$  features selected at random

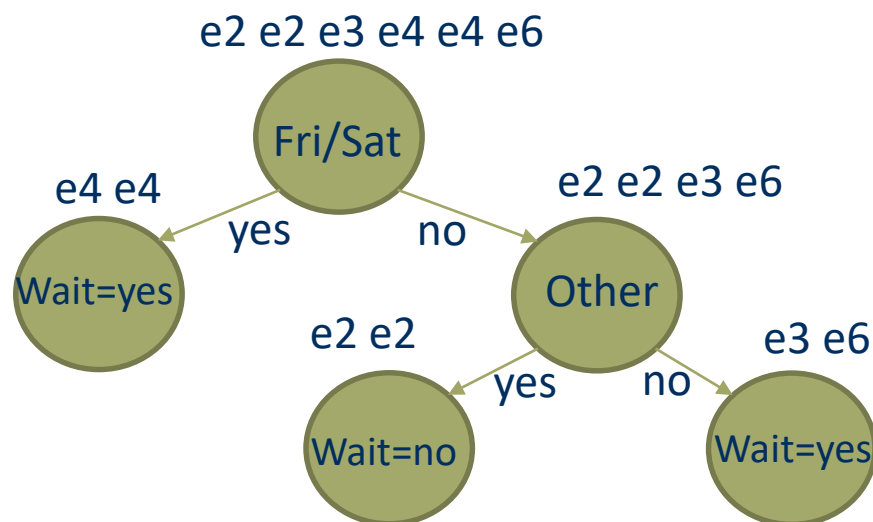
^T. K. Ho. 1998. The random subspace method for constructing decision forests, IEEE Transactions on Pattern Analysis and Machine Intelligence, 20(8):832-844

# Feature selection example

```
>>> from sklearn.utils import resample
>>> import numpy as np
>>> e = np.array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
>>> no_features = int(np.sqrt(len(e)))
>>> samp = resample(e, n_samples=no_features)
>>> samp
array([7, 6, 3])
>>> samp = resample(e, n_samples=no_features)
>>> samp
array([9, 3, 1])
>>> samp = resample(e, n_samples=no_features)
>>> samp
array([8, 0, 5])
>>> samp = resample(e, n_samples=no_features)
>>> samp
array([9, 2, 4])
>>> samp = resample(e, n_samples=no_features)
>>> samp
array([8, 3, 1])
```


# Feature selection example

Ex.	Other	Bar	Fri/Sat	Hungry	Guests	Wait
e2	yes	no	no	yes	full	no
e2	yes	no	no	yes	full	no
e3	no	yes	no	no	some	yes
e4	yes	no	yes	yes	full	yes
e4	yes	no	yes	yes	full	yes
e6	no	yes	no	yes	full	yes




# Random Forest

- Construct decision trees on bootstrap replicas
  - **Restrict** the node **decisions** to a small **subset** of **features** picked **randomly** for each node
- Do not prune the trees
  - Estimate tree performance on **out-of-bootstrap** data
  - Approximately **35%** of data in a bootstrap sample
  - Predict new instance
    - by majority vote (classification)
    - by averaging value (regression)





# Random forest example result



Breast-cancer-Wisconsin dataset from UCI repository  
Stratified 10 fold cross validation  
Decision trees as base classifier

## Random forest tips

- In (Fernandez-Delgado et al 2014), an **empirical** investigation was presented using:
  - 179 classifiers from 17 families, incl. all standard approaches
  - 121 datasets, incl. all of UCI except the largest ones
  - The random forest versions **ranked** the highest and they obtained near to the best accuracy for almost all the data sets
- It was **concluded** that
  - The classifiers most likely to be the best are the random forest (RF) versions

M. Fernandez-Delgado, E. Cernadas, S. Barro, and D. Amorim. Do we Need Hundreds of Classifiers to Solve Real World Classification Problems? *Journal of Machine Learning Research*, 15(1), pp. 3133-3181, 2014.

# Randomization


- Repeated **application** of a base learner that contains some **stochastic** element (or that can be adjusted to **contain** stochastic elements) such as
  - the setting of **initial** weights in a neural network
  - **sampling** of grow and prune data
  - choice of **attribute** to split on in decision trees
- Randomization may be **combined** with other techniques

# Boosting - AdaBoost

- **Adaptive Boosting** -> AdaBoost
- Idea: Turn a **weak stable** learner into a **strong** learner by focusing on the difficult cases to predict.
- Each predictor is created by using a **biased** sample of the **training** data, hence boosting minimize the bias of the ensemble (and also variance).
- General Steps in boosting
  1. Train a **weak** model on some training data
  2. Compute the **error** of the model on each training example
  3. Give **higher** importance to **examples** on which the model made **mistakes**
  4. Re-train the model using **re-weighted** training example


# AdaBoosting

- An **sequential** procedure to **adaptively change** the **distribution** of training data by
  - Focusing more on previously **miss-classified** examples
    - Initially all examples are assigned **equal** weights
    - After a classifier  $C_i$  is learned
      - The weights are **adjusted** to **increase** weights of **misclassified** examples of  $C_i$
      - Using **accuracy** on a validation set as the **basis** for **re-weighting**



Y. Freund and R. Schapire. 1997. A Decision-Theoretic Generalization of On-Line Learning and an Application to Boosting. J. Comput. Syst. Sci. 55(1): 119-139

# AdaBoost



- **Weak stable** learner – decision stump

Input: instances  $e_1, \dots, e_m$ , base learner BL, size  $S$

Output: a set of model-weight pairs  $M$

$w_1, \dots, w_m = 1$

$M = \{\}$

for  $i = 1$  to  $S$ :

$M_i = \text{BL}(\{(e_1, w_1), \dots, (e_m, w_m)\})$

$\text{Err} = (w_1 \cdot \text{err}(M_i, e_1) + \dots + w_m \cdot \text{err}(M_i, e_m)) / (w_1 + \dots + w_m)$

if  $\text{Err} = 0$  or  $\text{Err} > 0.5$  then break

$\alpha_i = \frac{1}{2} \ln((1 - \text{Err}) / \text{Err})$

for  $j = 1$  to  $m$ :



if  $\text{err}(M_i, e_j) = 0$  then  $w_j = w_j \cdot e(-\alpha_i)$  #down-w

else  $w_j = w_j \cdot e(\alpha_i)$  #up-weighted

Normalize( $w_1, \dots, w_m$ )

$M = M + (M_i, \alpha_i)$

# AdaBoost - Example



Err = 0.3      M1

$$\alpha_i = \frac{1}{2} \ln\left(\frac{1-\text{Err}}{\text{Err}}\right) = \frac{1}{2} \ln\left(\frac{1-0.3}{0.3}\right) = \mathbf{0.42}$$

$$w_j = w_j * e(-0.42) = 0.065$$

#down-weighted

$$w_j = w_j * e(0.42) = 0.152$$

#up-weighted

Normalize

$$0.065 * 7 = 0.455$$

$$0.152 * 3 = 0.456$$

$$0.455 + 0.456 = 0.911$$


$$0.455/0.911 = 0.5$$

$$0.5/7 = \mathbf{0.0714}$$

$$0.456/0.911 = 0.5$$

$$0.5/3 = \mathbf{0.167}$$

# AdaBoost - Example



$$\text{Err} = 0.0714 * 3 = 0.21 \quad \text{M2}$$

$$\alpha_i = \frac{1}{2} \ln\left(\frac{1-\text{Err}}{\text{Err}}\right) = \frac{1}{2} \ln\left(\frac{1-0.21}{0.21}\right) = 0.65$$

$$w_j = w_j * e^{-0.65} = 0.0714 * e^{-0.65} = 0.37$$

$$w_j = w_j * e^{-0.65} = 0.167 * e^{-0.65} = 0.87$$

$$w_j = 0.0714 * e^{0.65} = 1.37$$

$$0.37 * 4 = 1.48$$

$$0.87 * 3 = 2.61$$

$$1.37 * 3 = 4.11$$

$$1.48 + 2.61 + 4.11 = 8.2$$


$$1.48/8.2 = 0.18/4 = .045$$

$$2.61/8.2 = 0.32/3 = .106$$

$$4.11/8.2 = 0.5/3 = .166$$



# AdaBoost - Example




$$\text{Err} = 0.045 * 3 = \mathbf{0.14}$$

$$\alpha_i = \frac{1}{2} \ln\left(\frac{1-\text{Err}}{\text{Err}}\right) = \frac{1}{2} \ln\left(\frac{1-0.14}{0.14}\right) = \mathbf{0.92}$$



- Suppose we **stop** induction at round 3
- Then the ensemble consist of 3 predictors:
  - M1, M2, M3

# AdaBoost - Example




- Final predictor is a weighted linear combination of all predictors
- Where each predictor has weight  $\alpha_i$
- Hence **multiple** weak, linear classifiers **combined** to give a strong nonlinear predictor

# Decision stump



# AdaBoost example result



Breast-cancer-Wisconsin dataset from UCI repository  
Stratified 10 fold cross validation  
Decision stump as base classifier

# Gradient Boosting Machine - Regression

- Generalize Boosting methodology to use
  - Any **differentiable** loss functions
  - Gradient for mean squared error

<i>Setting</i>	<i>Loss function</i>	$-\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)}$
<i>Regression</i>	$\frac{1}{2} [y_i - f(x_i)]^2$	$y_i - f(x_i)$

Input: instances  $(y_1, x_1), \dots, (y_m, x_m)$ , base learner BL, size  $s$ , learning rate  $l$

Output: a sequence of models  $M_0, \dots, M_s$

$M_0 = (y_1 + \dots + y_m)/m$

#step 1 Initial model

for  $i = 1$  to  $s$ :

  for  $j = 1$  to  $m$ :

$y_j = (y_j - M_{i-1}(x_j))$

#step 2 Calculate residual values

$M_i = M_{i-1} + l * BL(\{(y_1, x_1), \dots, (y_m, x_m)\})$

#step 3 Induce new tree + update  $M_i$

# Gradient Boosting Machine - Regression

Input: instances  $(y_1, x_1), \dots, (y_m, x_m)$ , base learner BL, size  $s$ , learning rate  $\lambda$

Output: a sequence of models  $M_0, \dots, M_s$

$$M_0 = (y_1 + \dots + y_m)/m$$

#step 1 Initial model

for  $i = 1$  to  $s$ :

for  $j = 1$  to  $m$ :

$$r_j = (y_j - M_{i-1}(x_j))$$

#step 2 Calculate residual values

$$M_i = M_{i-1} + \lambda * BL(\{(y_1, x_1), \dots, (y_m, x_m)\})$$

#step 3 Induce new tree + update  $M_i$

Weight	Hair	Length
65	Short	178
55	Long	160
78	Long	180
95	Short	193

Step 1 Initial model:

$$M_0 = (y_1 + \dots + y_m)/m$$

$$M_0 = (178 + 160 + 180 + 193)/4 = 177.75$$

Step 2 Calculate residual values :

$$r_1 = (y_1 - M_{1-1}(x_1)) = (178 - 177.75) = 0.25$$

$$r_2 = (y_2 - M_{2-1}(x_1)) = (160 - 177.75) = -17.75$$

$$r_3 = (180 - 177.75) = 2.25$$

$$r_4 = (193 - 177.75) = 15.25$$

# Gradient Boosting Machine - Regression

Input: instances  $(y_1, x_1), \dots, (y_m, x_m)$ , base learner BL, size  $s$ , learning rate  $\eta$

Output: a sequence of models  $M_0, \dots, M_s$

$M_0 = (y_1 + \dots + y_m) / m$

#step 1 Initial model

for  $i = 1$  to  $s$ :

for  $j = 1$  to  $m$ :


$y_j = (y_j - M_{i-1}(x_j))$

#step 2 Calculate residual values

$M_i = M_{i-1} + \eta * BL(\{(y_1, x_1), \dots, (y_m, x_m)\})$  #step 3 Induce new tree + update  $M_i$

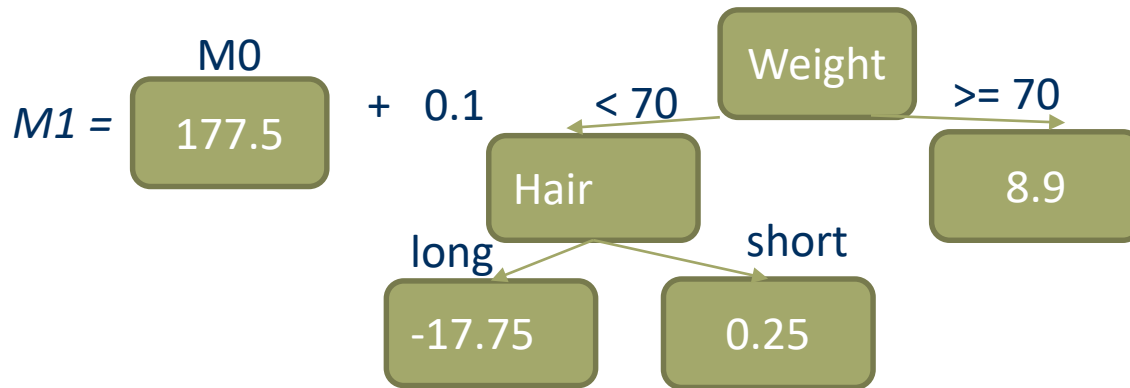
Step 3 Induce model  $M_1$ :

Weight	Hair	Length	Residual
65	Short	178	0.25
55	Long	160	-17.75
78	Long	180	2.25
95	Short	193	15.25



# Gradient Boosting Machine - Regression

- Predict regression value on the training data (learning rate[0,1] = 0.1):



Weight	Hair	Length	Residual	N. Res
65	Short	178	0.25	0.2
55	Long	160	-17.75	-15.9
78	Long	180	2.25	1.6
95	Short	193	15.25	14.6

Calculate new residual values:

$$Y1 = 177.75 + 0.1 * 0.25 = 177.8$$

$$Y2 = 177.75 + 0.1 * -17.75 = 175.9$$

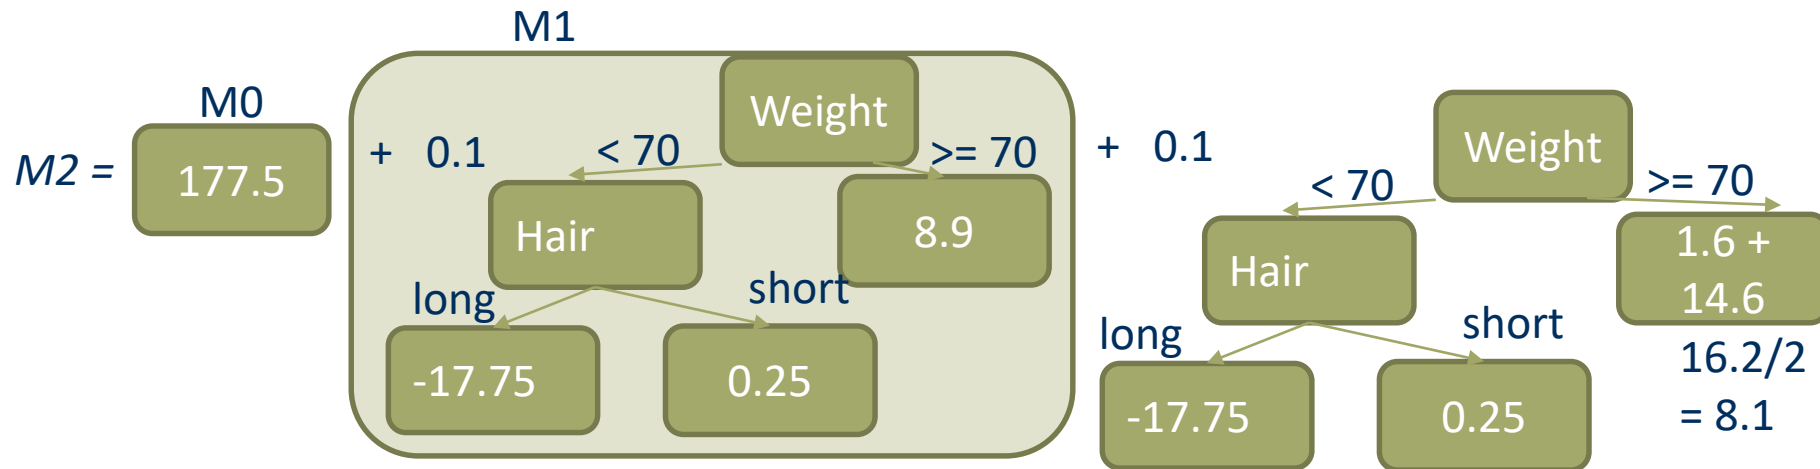
$$Y3 = 177.75 + 0.1 * 8.9 = 178.4$$

$$Y4 = 177.75 + 0.1 * 8.9 = 178.4$$



# Gradient Boosting Machine - Regression

Induce new tree:



Weight	Hair	Length	Residual	N. res.
65	Short	178	0.2	0.18
55	Long	160	-15.9	-15.5
78	Long	180	1.6	0.8
95	Short	193	14.6	13.8

Calculate new residual values:

$$Y_1 = 177.8 + 0.1 \cdot 0.2 = 177.82$$


$$Y_2 = 175.9 + 0.1 \cdot -15.9 = 174.3$$

$$Y_3 = 178.4 + 0.1 \cdot 8.1 = 179.2$$

$$Y_4 = 178.4 + 0.1 \cdot 8.1 = 179.2$$

# Gradient Boosting Machine - Regression

Using the model:




Weight	Hair	Length
80	long	?

$$\text{Predict} = 177.75 + 0.1 * 8.9 + 0.1 * 8.1 = 179.45$$

# Boosting tips

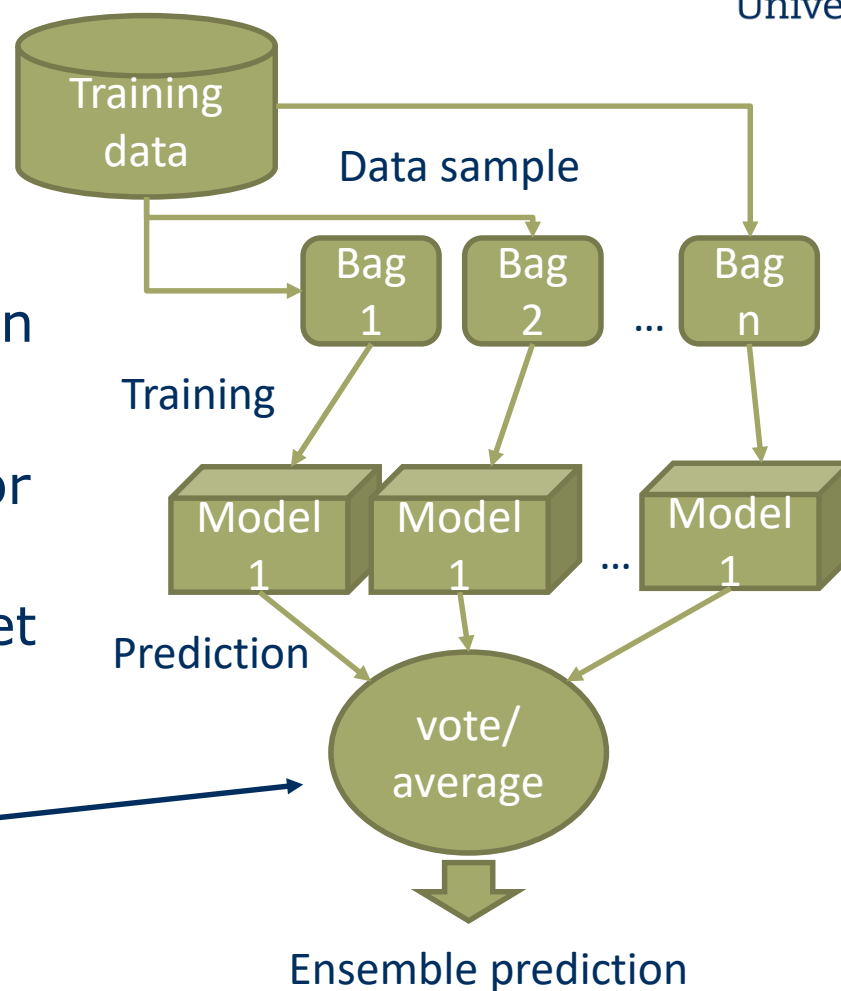
- For **multi-class** problems, it may be **difficult** for AdaBoost to reduce the error below 50% with weak base learners (like decision stumps)
  - More **powerful** base learners may be employed
  - The problem may be transformed into **multiple** binary classification problems (one against all)
  - Specific multi-class versions of AdaBoost have been developed
- Various **loss functions** may be used for the Gradient Boosting Machine (GBM), including log likelihood for classification
- GBM requires **tuning**
  - no. of iterations, model size and learning rate
- **XGBoost** implements GBM
  - well-known for winning many competitions, see ([www.kaggle.com](http://www.kaggle.com))
- Boosting can be **sensitive** to noise
  - erroneously labeled training examples

# Gradient Boosting example result



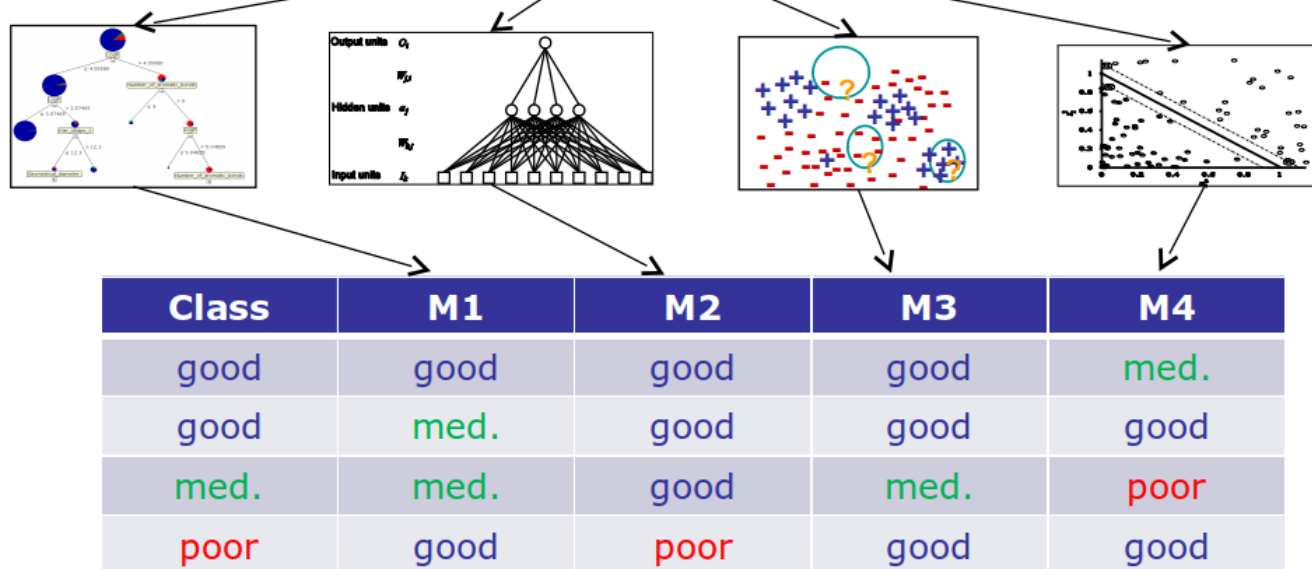
# Stacking

- Note that for example voting/averaging can be done in a weighted fashion
  - where weight come from for example
    - Accuracy on validation set
    - Or RMSError
- One can also use another **classifier** here
  - this is then called **stacking**



# Stacking Example

Name	Solubility	Frac. of rotatable bonds	Geom. diam.	Log P	Zagreb group index	Topol. diam.	No. C atoms	No. heavy bonds
methylopentane	good	0.40	3.40	2.44	19	4	6	5
methylocyclohexene	good	0	3.00	2.51	31	4	7	7
nonene	med.	0.75	6.93	3.53	28	8	9	8
hexadiene	good	0.60	4.30	2.14	16	5	6	5
butadiene	good	0.33	2.65	1.30	8	3	4	3
naphthalene	good	0	3.61	2.84	57	5	10	11
acenaphthylene	good	0	3.58	3.32	83	5	12	14
pyrene	poor	0	3.00	4.58	117	7	16	19
dimethylanthracene	poor	0	3.29	4.61	108	7	16	18
hexahydronaphthalene	med.	0	3.60	3.62	117	7	16	19
tdphenylene	poor	0	3.00	5.15	126	7	18	21
benzo[c]pyrene	poor	0	3.29	5.64	152	7	20	24





IF M2 = good & M3 = good THEN Class = good  
 IF M2 = poor THEN Class = poor

...

# Stacking tips

- **Linear models** have been shown to be effective when learning the **combination** function
- Predictive performance can be improved by combining class **probability estimates** rather than **class labels** generated by the base models
- Why is this?

# Interpretability of Ensembles



- There exist different approaches to estimating variable importance
  - one method is to measure the relative performance degradation (on OOB predictions) when permuting the values of each feature in turn.



# Summary

- Bias/Variance **tradeoff**
  - Ensemble methods that minimize **variance**
    - Bagging
    - Random Forests
  - Ensemble methods that minimize **bias**
    - AdaBoost
    - Gradient Boosting Machine
- Compared to the base learning algorithms, ensembles
  - typically substantially **improve** the predictive performance
    - sometimes to **state-of-the-art** performance
- The increased predictive **power** comes
  - at the cost of **reduced** interpretability
- The combination strategies may can lead to **increased** computational time (both in learning and prediction)
  - But this time penalty can be **eliminated** via parallelization

# All for today

Questions?